

File Systems

1. NAME

dirent.h — format of directory entries

SYNOPSIS

```
#include <dirent.h>
```

DESCRIPTION

The internal format of directories is unspecified. The *<dirent.h>* header shall define the following type:

DIR A type representing a directory stream. The **DIR** type may be an incomplete type. It shall also define the structure **dirent** which shall include the following members:

ino_t d_ino File serial number.

char d_name [] Filename string of entry.

The array *d_name* is of unspecified size, but shall contain a filename of at most {NAME_MAX} bytes followed by a terminating null byte.

```
struct dirent {
    ino_t      d_ino;           /* inode number */
    off_t      d_off;          /* offset to the next dirent */
    unsigned short d_reclen;    /* length of this record */
    unsigned char d_type;       /* type of file;
char    d_name[256]; /* filename */ (This we have used in our ls-l Program)
}
```

2. NAME

sys/stat.h - data returned by the stat() function

SYNOPSIS

```
#include <sys/stat.h>
```

Description

The `<sys/stat.h>` header shall define the structure of the data returned by the functions [*fstat\(\)*](#), [*lstat\(\)*](#), and [*stat\(\)*](#).

The **stat** structure shall contain at least the following members:

- `mode_t st_mode` Mode of file (see below).
- `nlink_t st_nlink` Number of hard links to the file.
- `uid_t st_uid` User ID of file.
- `gid_t st_gid` Group ID of file.
- `off_t st_size` For regular files, the file size in bytes.
 - For symbolic links, the length in bytes of the
 - pathname contained in the symbolic link.
- `time_t st_atime` Time of last access.
- `time_t st_mtime` Time of last data modification.
- `time_t st_ctime` Time of last status change.
- `blksize_t st_blksize` A file system-specific preferred I/O block size for
 - this object. In some file system types, this may
 - vary from file to file.
- `blkcnt_t st_blocks` Number of blocks allocated for this object.

- The following macros shall be provided to test whether a file is of the specified type. The value *m* supplied to the macros is the value of *st_mode* from a **stat** structure. The macro shall evaluate to a non-zero value if the test is true; 0 if the test is false.
- **S_ISBLK(*m*)**
- **Test for a block special file.**
- **S_ISCHR(*m*)**
- **Test for a character special file.**
- **S_ISDIR(*m*)**
- **Test for a directory.**
- **S_ISFIFO(*m*)**
- **Test for a pipe or FIFO special file.**
- **S_ISREG(*m*)**
- **Test for a regular file.**
- **S_ISLNK(*m*)**
- **Test for a symbolic link.**
- **S_ISSOCK(*m*)**
- **Test for a socket.**

3. NAME

pwd.h - password structure

SYNOPSIS

#include <pwd.h>

DESCRIPTION

The *<pwd.h>* header provides a definition for **struct passwd**, which includes at least the following members:

char *pw_name user's login name

uid_t pw_uid numerical user ID

gid_t pw_gid numerical group ID

char *pw_dir initial working directory

char *pw_shell program to use as shell

The following are declared as functions and may also be defined as macros. Function prototypes must be provided for use with an ISO C compiler.

- **struct passwd *getpwnam(const char *);**
- **struct passwd *getpwuid(uid_t);**
- **int getpwnam_r(const char *, struct passwd *, char *,
 size_t, struct passwd **);**
- **int getpwuid_r(uid_t, struct passwd *, char *,
 size_t, struct passwd **);**
- **void endpwent(void);**
- **struct passwd *getpwent(void);**
-

4. NAME

grp.h - group structure

SYNOPSIS

#include <grp.h>

DESCRIPTION

The *<grp.h>* header declares the structure **group** which includes the following members:

- char *gr_name the name of the group
- gid_t gr_gid numerical group ID
- char **gr_mem pointer to a null-terminated array of character

The following are declared as functions and may also be defined as macros. Function prototypes must be provided for use with an ISO C compiler.

- struct group *getgrgid(gid_t);
- struct group *getgrnam(const char *);
- int getgrgid_r(gid_t, struct group *, char *,
size_t, struct group **);
- int getgrnam_r(const char *, struct group *, char *,
size_t, struct group **);
- struct group *getgrent(void);

5. NAME

time.h - time types

SYNOPSIS

```
#include <time.h>
```

The following are declared as functions and may also be defined as macros

- struct tm *[localtime](#)(const time_t *);
- struct tm *[localtime_r](#)(const time_t *, struct tm *);
- time_t [mktime](#)(struct tm *);
- int [nanosleep](#)(const struct timespec *, struct timespec *);
- size_t [strftime](#)(char *, size_t, const char *, const struct tm *);
- char *[strptime](#)(const char *, const char *, struct tm *);

An **Arrow operator** in C/C++ allows to access elements in [Structures](#) and [Unions](#). It is used with a [pointer variable pointing to a structure or union](#). The arrow operator is formed by using a minus sign, followed by the greater than symbol as shown below.

Syntax:

```
(pointer_name)->(variable_name)
```

```
de->d_name
```

CREATE SYSTEM CALL

```
#include <fcntl.h>
```

```
int creat(char* filename, mode_t mode)
```

- The mode
 - is an octal number
 - Example: 0444 indicates that r access for USER, GROUP and ALL for the file.
 - If the file exists, the creat is ignored and prior content and rights are maintained.

Opening Files

```
#include <sys/types.h>
```

```
#include <sys/stat.h>
```

```
#include <fcntl.h>
```

```
int open(char* filename, int flags, mode_t mode);
```

Flags: O_RDONLY, O_WRONLY, O_RDWR, O_CREAT, O_TRUNC, O_APPEND

– Mode: Specifies permission bits of the file

- S_IRUSR, S_IWUSR, S_IXUSR – owner permission
- S_IRGRP, S_IWGRP, S_IXGRP – group permission
- S_IROTH, S_IWOTH, S_IXOTH – other permission

Reading/Writing Files

- Low level read and write
- `#include <unistd.h>`
- `ssize_t read(int fd, void *buf, size_t n);`

Returns num bytes read or -1

- `ssize_t write(int fd, const void *buf, size_t n);`
 - Returns num bytes written or -1

What about `size_t` and `ssize_t`

- `size_t` – unsigned int
- `ssize_t` - signed int
- How does this affect the range of values in each type?
 - with 32-bit int?

Reading file metadata

- How can we find information about a file
- `#include <unistd.h>`
- `#include <sys/stat.h>`
`int stat(const char* filename, struct stat *buf);`
- `int fstat(int fd, struct stat *buf);`

What is struct stat?

```
struct stat
{
    dev_t      st_dev;      /* ID of device containing file */
    ino_t      st_ino;      /* inode number */
    mode_t     st_mode;     /* protection */
    nlink_t    st_nlink;    /* number of hard links */
    uid_t      st_uid;      /* user ID of owner */
    gid_t      st_gid;      /* group ID of owner */
    dev_t      st_rdev;     /* device ID (if special file) */
    off_t      st_size;     /* total size, in bytes */
    blksize_t  st_blksize;  /* blocksize for filesystem I/O */
    blkcnt_t   st_blocks;   /* number of blocks allocated */
    time_t     st_atime;    /* time of last access */
    time_t     st_mtime;    /* time of last modification */
    time_t     st_ctime;    /* time of last status change */
};
```

Accessing File Status

`stat(char* file, struct stat *buf);`

`fstat(int fd, struct stat *buf);`

`struct stat buf; // defines a struct stat to hold file information`

`stat("filename", &buf) ; // now the file information is placed in the buf`

`st_atime` --- Last access time

`st_mtime` --- last modify time

`st_ctime` --- Last status change time

`st_size` --- total size of file

`st_uid` – user ID of owner

`st_mode` – file status (directory or not)

Example

```
#include <sys/types.h>
#include <sys/stat.h>
#include <dirent.h>
struct stat statbuf;

char dirpath[256];
getcwd(dirpath,256);

DIR *dir = opendir(dirpath);
struct dirent *dp;

for (dp=readdir(dir); dp != NULL ; dp=readdir(dir)){
    stat(dp->d_name, &statbuf);

    printf("the file name is %s \n", dp->d_name);
    printf("dir = %d\n", S_ISDIR(statbuf.st_mode));
    printf("file size is %ld in bytes \n", statbuf.st_size);

    printf("last modified time is %ld in seconds \n", statbuf.st_mtime);
    printf("last access time is %ld in seconds \n", statbuf.st_atime);
    printf("The device containing the file is %d\n", statbuf.st_dev);
    printf("File serial number is %d\n\n", statbuf.st_ino);
}
```

How to determine a file type

- S_ISREG
 - A regular file?
- S_ISDIR
 - Is a directory?
 - `printf("dir = %d\n", S_ISDIR(statbuf.st_mode));`
- S_ISSOCK

A network socket

The Breakdown of LS-L Program

DIR *the directory

The directory: it's the folder we're browsing (we'll use an argument (argv) in order to identify it)

struct dirent *de

A dirent structure contains the character pointer d_name, which points to a string that gives the name of a file in the directory. This string ends in a terminating NULL, and has a maximum of NAME_MAX characters.

struct stat buf;

// defines a struct stat to hold file information

buf

(Output) A pointer to the area to which the information should be written.

struct passwd *p;

struct group *p;

Will be used to determine the file owner & group

The **opendir()** function opens a directory stream corresponding to the directory name, and returns a pointer to the directory stream. The stream is positioned at the first entry in the directory.

If a file is found (readdir returns a NOT NULL value), the loop starts/keep going until it has listed all of them.

```
while( (de=readdir(d)) !=NULL)
```

Then we use stat function in order to retrieve information about the file

```
stat(buf, &thestat);
```

Then we will find the file permissions using the loop and group name, user id and time using the respective system calls

User Name

```
p=getpwuid(buf.st_uid);  
printf(" %.8s",p->pw_name);
```

Group Name

```
g=getgrgid(buf.st_gid);  
printf(" %-8.8s",g->gr_name);
```


For Time

The C library function `struct tm *localtime(const time_t *timer)` uses the time pointed by `timer` to fill a `tm` structure with the values that represent the corresponding local time. The value of `timer` is broken up into the structure `tm` and expressed in the local time zone.

Declaration

Following is the declaration for `localtime()` function.

```
struct tm *localtime(const time_t *timer)
```

`strftime()` is a function in C which is used to format date and time. It comes under the header file `time.h`, which also contains a structure named `struct tm` which is used to hold the time and date. The syntax of `strftime()` is as shown below

```
size_t strftime(char *s, size_t max, const char *format, const struct tm *tm);
```